

▼ Siddharth Patel

19BCP130

```
import numpy as np
import pandas as pd

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

import matplotlib.pyplot as plt
```

```
dataset = pd.read_csv("./train.csv")
X = dataset.drop(['label'], axis=1).values
y = dataset['label'].values
dataset
```



| | label | pixel0 | pixel1 | pixel2 | pixel3 | pixel4 | pixel5 | pixel6 | pixel7 | pixel8 | ... | pixel |
|-------|-------|--------|--------|--------|--------|--------|--------|--------|--------|--------|-----|-------|
| 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |
| 2 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |
| 3 | 4 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |
| 4 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 41995 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |
| 41996 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |
| 41997 | 7 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |
| 41998 | 6 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |
| 41999 | 9 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | |

42000 rows × 785 columns

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.fit_transform(X_test)
```

```
X_train_scaled
```

```
array([[0., 0., 0., ..., 0., 0., 0.],
```

```
[0., 0., 0., ..., 0., 0., 0.],
[0., 0., 0., ..., 0., 0., 0.],
...,
[0., 0., 0., ..., 0., 0., 0.],
[0., 0., 0., ..., 0., 0., 0.],
[0., 0., 0., ..., 0., 0., 0.]])
```

```
import warnings
warnings.filterwarnings("ignore")

parameters = {'C': [.001, .01, .1, 1, 10, 100]}

grid = GridSearchCV(estimator=LogisticRegression(), param_grid=parameters, cv=4, verbose=2, return_t
grid.fit(X_train_scaled, y_train)
```

Fitting 4 folds for each of 6 candidates, totalling 24 fits

```
[CV] END .....C=0.001; total time= 11.4s
[CV] END .....C=0.001; total time= 11.7s
[CV] END .....C=0.001; total time= 11.5s
[CV] END .....C=0.001; total time= 15.1s
[CV] END .....C=0.01; total time= 9.9s
[CV] END .....C=0.01; total time= 9.7s
[CV] END .....C=0.01; total time= 9.8s
[CV] END .....C=0.01; total time= 9.6s
[CV] END .....C=0.1; total time= 9.7s
[CV] END .....C=0.1; total time= 9.6s
[CV] END .....C=0.1; total time= 9.4s
[CV] END .....C=0.1; total time= 9.6s
[CV] END .....C=1; total time= 9.9s
[CV] END .....C=1; total time= 9.5s
[CV] END .....C=1; total time= 9.6s
[CV] END .....C=1; total time= 9.5s
[CV] END .....C=10; total time= 9.6s
[CV] END .....C=10; total time= 9.5s
[CV] END .....C=10; total time= 9.3s
[CV] END .....C=10; total time= 9.3s
[CV] END .....C=100; total time= 9.4s
[CV] END .....C=100; total time= 9.6s
[CV] END .....C=100; total time= 9.5s
[CV] END .....C=100; total time= 9.6s
```

```
GridSearchCV(cv=4, estimator=LogisticRegression(),
              param_grid={'C': [0.001, 0.01, 0.1, 1, 10, 100]},
              return_train_score=True, verbose=2)
```

```
grid.cv_results_
```

```
{'mean_fit_time': array([12.47271502, 9.81256151, 9.6483801 , 9.69212663, 9.49873281,
9.59169298]),
 'std_fit_time': array([1.58217727, 0.13810486, 0.10798535, 0.15484578, 0.09669219,
0.05857961]),
 'mean_score_time': array([0.03489959, 0.02898633, 0.02822733, 0.02880573, 0.02705884,
0.0275991 ]),
 'std_score_time': array([0.00275724, 0.00388277, 0.0005116 , 0.00140212, 0.00077676,
0.00078495]),
 'param_C': masked_array(data=[0.001, 0.01, 0.1, 1, 10, 100],
                          mask=[False, False, False, False, False, False],
                          fill_value='?',
                          dtype=object),
 'params': [{'C': 0.001},
 {'C': 0.01},
 {'C': 0.1},
 {'C': 1},
```

```

{'C': 10},
{'C': 100}],
'split0_test_score': array([0.9125      , 0.92202381, 0.915      , 0.90154762, 0.89547619,
0.89559524]),
'split1_test_score': array([0.90797619, 0.91642857, 0.90928571, 0.89738095, 0.89261905,
0.89011905]),
'split2_test_score': array([0.91083333, 0.91964286, 0.91404762, 0.89988095, 0.89428571,
0.89297619]),
'split3_test_score': array([0.91511905, 0.92083333, 0.91357143, 0.90011905, 0.89464286,
0.89297619]),
'mean_test_score': array([0.91160714, 0.91973214, 0.91297619, 0.89973214, 0.89425595,
0.89291667]),
'std_test_score': array([0.0025939 , 0.00208482, 0.0021919 , 0.00149966, 0.00103911,
0.00193704]),
'rank_test_score': array([3, 1, 2, 4, 5, 6]),
'split0_train_score': array([0.92043651, 0.94099206, 0.95535714, 0.96253968, 0.96384921,
0.96353175]),
'split1_train_score': array([0.92186508, 0.94293651, 0.95805556, 0.96488095, 0.96619048,
0.96686508]),
'split2_train_score': array([0.92103175, 0.94305556, 0.95746032, 0.96563492, 0.96607143,
0.96626984]),
'split3_train_score': array([0.91964286, 0.94170635, 0.95619048, 0.96436508, 0.965      ,
0.96519841]),
'mean_train_score': array([0.92074405, 0.94217262, 0.95676587, 0.96435516, 0.96527778,
0.96546627]),
'std_train_score': array([0.00081343, 0.0008623 , 0.00105607, 0.0011413 , 0.00094616,
0.00126654])})

```

```
C_vals = [str(x) for x in parameters['C']]
```

```


plt.xlabel("C")
plt.ylabel("cv-accuracy")
plt.scatter(C_vals, grid.cv_results_['mean_test_score'])
plt.vlines(C_vals, 0, grid.cv_results_['mean_test_score'], linestyle="dashed")
plt.ylim(0.80,.95)
plt.xticks(C_vals)
plt.show()

```

```

plt.xlabel("C")
plt.ylabel("train-accuracy")
plt.scatter(C_vals, grid.cv_results_['mean_train_score'])
plt.vlines(C_vals, 0, grid.cv_results_['mean_train_score'], linestyle="dashed")
plt.ylim(0.85,1.00)
plt.xticks(C_vals)
plt.show()

```



```
print("Best parameters:", grid.best_params_)
print("Best score:", grid.best_score_)
```

```
model = LogisticRegression(C=.01)
```

```
model.fit(X_train_scaled, y_train)
y_pred = model.predict(X_test_scaled)
```

```
accuracy_score(y_test, y_pred)
```

```
0.9195238095238095
```

```
for c in parameters['C']:
    model = LogisticRegression(C=c)

    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_test_scaled)

    print(c, accuracy_score(y_test, y_pred))
```

```
0.001 0.9125
0.01 0.9195238095238095
0.1 0.916547619047619
1 0.9089285714285714
10 0.9060714285714285
100 0.9058333333333334
```

